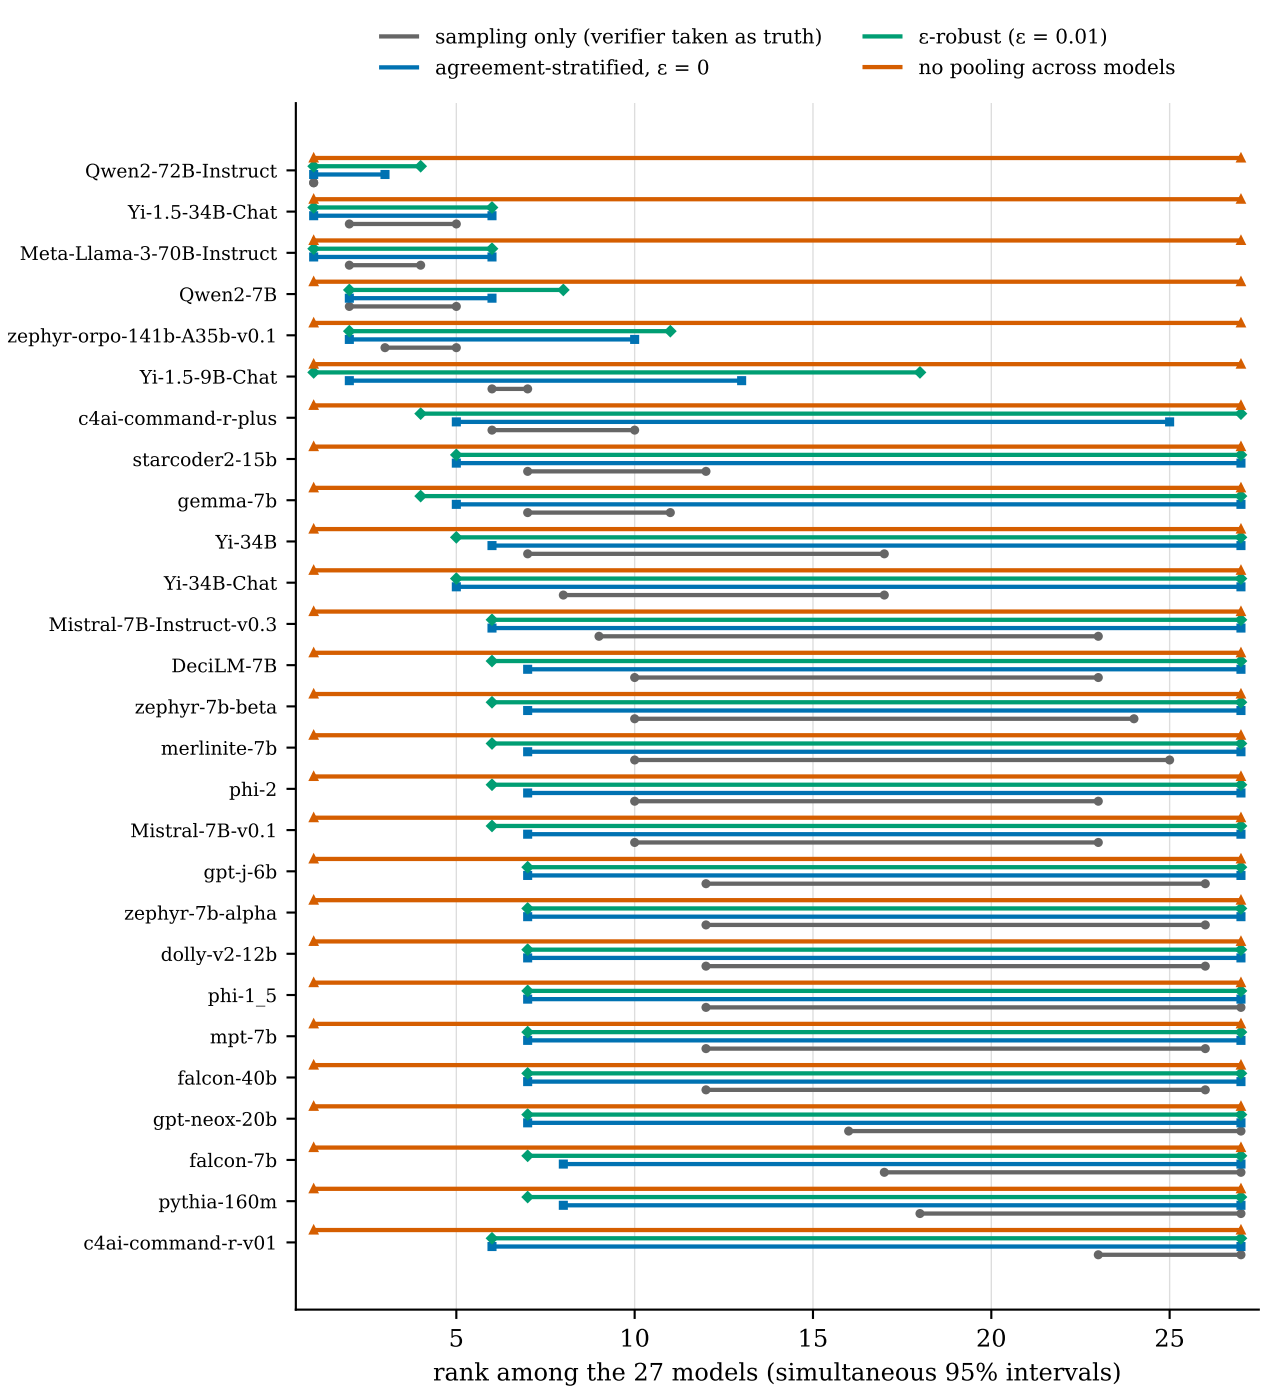
sampling only (verifier taken as truth)
agreement-stratified, ε = 0
ε-robust (ε = 0.01)
no pooling across models
Qwen2-72B-Instruct
Yi-1.5-34B-Chat
Meta-Llama-3-70B-Instruct
Qwen2-7B
zephyr-orpo-141b-A35b-v0.1
Yi-1.5-9B-Chat
c4ai-command-r-plus
starcoder2-15b
gemma-7b
Yi-34B
Yi-34B-Chat
Mistral-7B-Instruct-v0.3
DeciLM-7B
zephyr-7b-beta
merlinite-7b
phi-2
Mistral-7B-v0.1
gpt-j-6b
zephyr-7b-alpha
dolly-v2-12b
phi-1_5
mpt-7b
falcon-40b
gpt-neox-20b
falcon-7b
pythia-160m
c4ai-command-r-v01
5
10
15
20
25
rank among the 27 models (simultaneous 95% intervals)